

# Operating Systems–2: Spring 2024

## Programming Assignment 3:

### Dynamic Matrix Squaring

### Report

Prabhat - CO22BTECH11009

#### Program Overview:

Below is an overview of the program:

- ❑ In this programming assignment, we will build upon our previous work measuring the performance of calculating the square of matrix A in parallel in C++. Previously, we allocated the rows of the matrix to threads statically (in either mixed or chunk mode).
- ❑ In this assignment, the sequence of rows will be allocated dynamically to the threads.
- ❑ Each thread increments a shared counter C by a value of rowInc (row Increment) to claim the set of rows of the C matrix it is responsible for. The value rowInc can be seen similarly to the chunk size of assignment 1, where the chunk size was fixed statically.
- ❑ There will be synchronization issues when threads compete to increment the counter C. To solve this problem, I have implemented different mutual exclusion algorithms to increment C, which are: (a) TAS, (b) CAS, (c) Bounded CAS, and (d) atomic increment (provided by the C++ atomic library).

## Specifications of the System:

### ❑ Hardware Information:

1. Hardware Model: ASUSTeK COMPUTER INC. ASUS TUF Gaming A15 FA506ICB\_FA506ICB
2. Memory: 16.0 GiB
3. Processor: AMD Ryzen™ 7 4800H with Radeon™ Graphics × 16
4. Graphics: AMD Radeon™ Graphics
5. Disk Capacity: 512.1 GiB
6. CPU Cores: 8 (Octa core)
7. Threads: 16

### ❑ Software Information:

1. Firmware Version: FA506ICB.304
2. OS Name: Ubuntu 23.10
3. OS Build: (null)
4. OS Type: 64-bit
5. Windowing System: Wayland
6. Kernel Version: Linux 6.5.0-15-generic

## TAS Design:

This program performs matrix multiplication with Test-and-Set (TAS) synchronization. The low-level design for the chunk program is as follows:

### 1. Atomic Variables:

- ☐ `atomic int counter(0)`: An atomic integer counter initialized to 0, used to keep track of the current row being processed by threads.
- ☐ `atomic_flag lock = ATOMIC_FLAG_INIT`: An atomic flag 'lock' initialized to clear state, used for mutual exclusion to prevent multiple threads from accessing the critical section simultaneously.

### 2. Matrix Multiplication Function ('multiply'):

- ☐ This function is called by each thread to compute a portion of the resulting matrix.
- ☐ It takes references to matrices A (input matrix) and C (result matrix), the size of the matrix N, and the row increment rowInc.
- ☐ It uses a while loop to repeatedly acquire the lock using `test_and_set()` function to enter the critical section.
- ☐ Inside the critical section, it updates `start_row` and `end_row` based on the current value of counter, increments counter, and releases the lock.
- ☐ It then performs matrix multiplication for the assigned rows.
- ☐ The loop breaks when all rows have been processed.

### 3. Thread Assignment Function ('TAS'):

- ☐ This function creates and manages multiple threads to perform matrix multiplication concurrently.
- ☐ It takes references to matrices A and C, the size of the matrix N, the number of threads K, and the row increment rowInc.
- ☐ It creates K threads, each calling the 'multiply' function with appropriate parameters.
- ☐ After creating all threads, it waits for each thread to finish using `'join()'`

### 4. Main Function ('main'):

- ☐ Reads input parameters (N, K, rowInc) and matrix A from a file (inp.txt).
- ☐ Calls the TAS function to perform matrix multiplication.
- ☐ Measures the execution time using the chrono library.
- ☐ Writes the result matrix C and the execution time to an output file (out\_TAS.txt).

5. Input and Output:

- ❑ The input matrix A, size parameters N, K, rowInc are read from a file (inp.txt).
- ❑ The resulting matrix C and the execution time are written to an output file (out\_TAS.txt).

6. Synchronization:

- ❑ Atomic operations ('test\_and\_set()', 'clear()') are used for mutual exclusion to ensure that only one thread accesses the critical section at a time.

Overall, the program efficiently utilizes multiple threads to perform matrix multiplication concurrently while ensuring synchronization using atomic operations.

## CAS Design:

This program performs matrix multiplication with Compare-and-Swap (CAS) synchronization. The low-level design for the chunk program is as follows:

### 1. Atomic Variables:

- ❑ `atomic int counter(0)`: An atomic integer counter initialized to 0, used to keep track of the current row being processed by threads.
- ❑ `atomic bool lock(false)`: An atomic boolean variable initialized to false, used for mutual exclusion to prevent multiple threads from accessing the critical section simultaneously.

### 2. Matrix Multiplication Function ('multiply'):

- ❑ This function is called by each thread to compute a portion of the resulting matrix.
- ❑ It takes references to matrices A (input matrix) and C (result matrix), the size of the matrix N, and the row increment rowInc.
- ❑ It employs a while loop to repeatedly acquire the lock using CAS operation to enter the critical section.
- ❑ Inside the critical section, it updates the start row and end row based on the current value of the counter, increments counter, and releases the lock.
- ❑ It then performs matrix multiplication for the assigned rows.
- ❑ The loop breaks when all rows have been processed.

### 3. Thread Assignment Function ('CAS'):

- ❑ This function creates and manages multiple threads to perform matrix multiplication concurrently.
- ❑ It takes references to matrices A and C, the size of the matrix N, the number of threads K, and the row increment rowInc.
- ❑ It creates K threads, each calling the multiply function with appropriate parameters.
- ❑ After creating all threads, it waits for each thread to finish using 'join()'

### 4. Main Function ('main'):

- ❑ Reads input parameters (N, K, rowInc) and matrix A from a file (inp.txt).
- ❑ Calls the CAS function to perform matrix multiplication.
- ❑ Measures the execution time using the chrono library.
- ❑ Writes the result matrix C and the execution time to an output file (out\_CAS.txt).

5. Input and Output:

- ❑ The input matrix A, size parameters N, K, rowInc are read from a file (inp.txt).
- ❑ The resulting matrix C and the execution time are written to an output file (out\_CAS.txt).

6. Synchronization:

- ❑ CAS (Compare-and-Swap) operation is used for mutual exclusion to ensure that only one thread accesses the critical section at a time.

Overall, the program efficiently utilizes multiple threads to perform matrix multiplication concurrently while ensuring synchronization using CAS operation for locking.

## Bounded CAS Design:

This program performs matrix multiplication with Bounded Compare-and-Swap (CAS) synchronization. The low-level design for the chunk program is as follows:

### 1. Atomic Variables:

- ❑ `atomic int counter(0)`: An atomic integer counter initialized to 0, used to keep track of the current row being processed by threads.
- ❑ `atomic bool lock(false)`: An atomic boolean variable initialized to false, used for mutual exclusion to prevent multiple threads from accessing the critical section simultaneously.
- ❑ `const int MAX_RETRIES = 1000`: Constant defining the maximum number of retries for acquiring the lock.

### 2. Matrix Multiplication Function ('multiply'):

- ❑ This function is called by each thread to compute a portion of the resulting matrix.
- ❑ It takes references to matrices A (input matrix) and C (result matrix), the size of the matrix N, and the row increment rowInc.
- ❑ It employs a while loop to repeatedly acquire the lock using CAS operation to enter the critical section, with a bounded number of retries.
- ❑ Inside the critical section, it updates the start row and end row based on the current value of the counter, increments the counter, and releases the lock.
- ❑ It then performs matrix multiplication for the assigned rows.
- ❑ The loop breaks when all rows have been processed or the maximum number of retries is reached.

### 3. Thread Assignment Function ('BoundedCAS'):

- ❑ This function creates and manages multiple threads to perform matrix multiplication concurrently.
- ❑ It takes references to matrices A and C, the size of the matrix N, the number of threads K, and the row increment rowInc.
- ❑ It creates K threads, each calling the multiply function with appropriate parameters.
- ❑ After creating all threads, it waits for each thread to finish using 'join()'

4. Main Function ('main'):

- ☐ Reads input parameters (N, K, rowInc) and matrix A from a file (inp.txt).
- ☐ Calls the BoundedCAS function to perform matrix multiplication.
- ☐ Measures the execution time using the chrono library.
- ☐ Writes the result matrix C and the execution time to an output file (out\_BoundedCAS.txt).

5. Input and Output:

- ☐ The input matrix A, size parameters N, K, rowInc are read from a file (inp.txt).
- ☐ The resulting matrix C and the execution time are written to an output file (out\_BoundedCAS.txt).

6. Synchronization:

- ☐ Bounded CAS (Compare-and-Swap) operation is used for mutual exclusion to ensure that only one thread accesses the critical section at a time, with a maximum number of retries to prevent indefinite waiting.

Overall, the program efficiently utilizes multiple threads to perform matrix multiplication concurrently while ensuring synchronization using bounded CAS operation for locking.



## Atomic Design:

This program performs matrix multiplication using atomic increment synchronization. The low-level design for the chunk program is as follows:

### 1. Atomic Variables:

- ❑ `atomic int counter(0)`: An atomic integer counter initialized to 0, used to keep track of the current row being processed by threads.

### 2. Matrix Multiplication Function ('multiply'):

- ❑ This function is called by each thread to compute a portion of the resulting matrix.
- ❑ It takes references to matrices A (input matrix) and C (result matrix), the size of the matrix N, and the row increment rowInc.
- ❑ It fetches the start row atomically by adding rowInc to the counter using `atomic_fetch_add_explicit` with relaxed memory order.
- ❑ It employs a while loop to iteratively process rows until all rows have been processed.
- ❑ Inside the loop, it computes the end row based on the start row and rowInc.
- ❑ If the start row is less than N, it performs matrix multiplication for the assigned rows and updates the start row atomically.
- ❑ The loop breaks when all rows have been processed.

### 3. Thread Assignment Function ('Atomic'):

- ❑ This function creates and manages multiple threads to perform matrix multiplication concurrently.
- ❑ It takes references to matrices A and C, the size of the matrix N, the number of threads K, and the row increment rowInc.
- ❑ It creates K threads, each calling the 'multiply' function with appropriate parameters.
- ❑ After creating all threads, it waits for each thread to finish using 'join()'

### 4. Main Function ('main'):

- ❑ Reads input parameters (N, K, rowInc) and matrix A from a file (inp.txt).
- ❑ Calls the Atomic function to perform matrix multiplication.
- ❑ Measures the execution time using the chrono library.
- ❑ Writes the result matrix C and the execution time to an output file (out\_Atomic.txt).

5. Input and Output:

- ❑ The input matrix A, size parameters N, K, rowInc are read from a file (inp.txt).
- ❑ The resulting matrix C and the execution time are written to an output file (out\_Atomic.txt).

6. Synchronization:

- ❑ Atomic increment operation (`atomic_fetch_add_explicit`) is used to fetch and update the start row atomically, ensuring that each thread processes a unique set of rows without interference.

Overall, the program efficiently utilizes multiple threads to perform matrix multiplication concurrently while ensuring synchronization using atomic increment operation for updating the start row.

## Experiment 1:

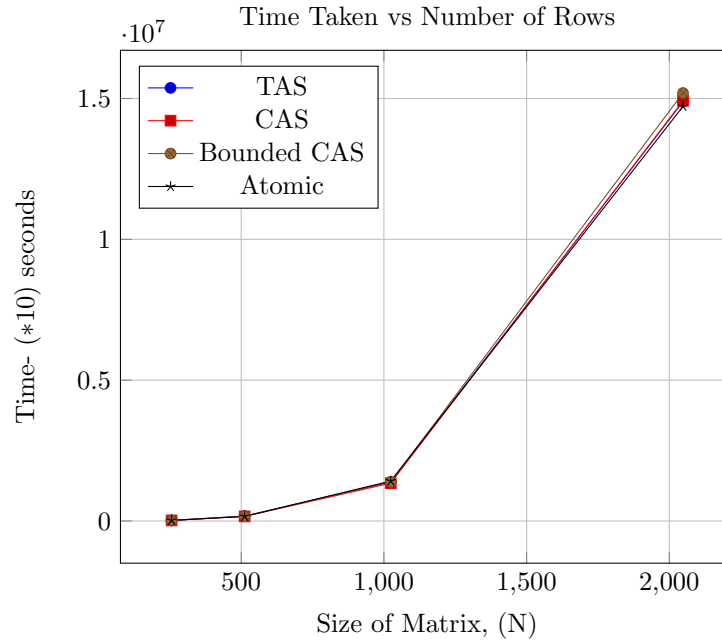


Figure 1: Time vs Size of Matrix (N)

### Observations:

- ❑ As the size of the matrix, N increases, the time taken to calculate the square of the matrix also increases.
- ❑ All the algorithms, i.e. TAS, CAS, Bounded CAS and Atomic methods takes almost the same time.
- ❑ Also the pattern in the curve as the N increases is almost same for all the curves.

## Experiment 2:

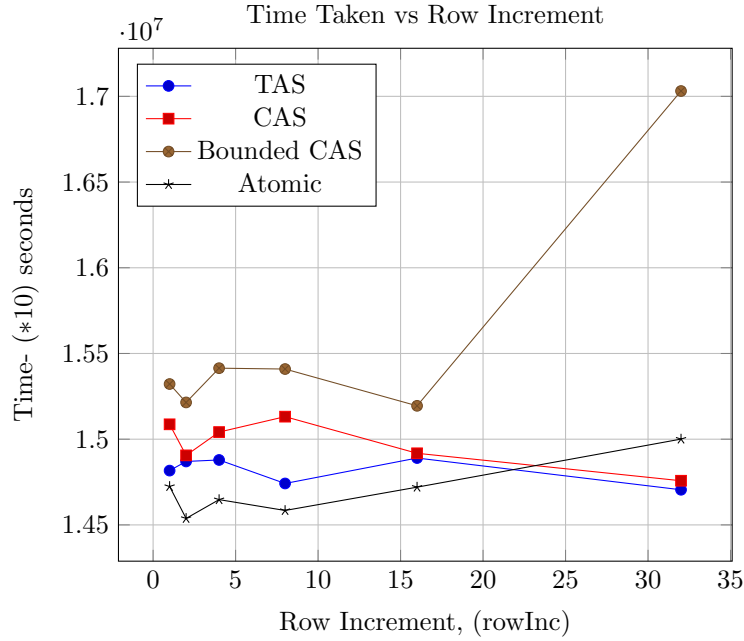


Figure 2: Time vs Row Increment (rowInc)

### Observations:

- ☐ There is no such pattern between the taken taken for different value rowInc.
- ☐ All the algorithms have almost the same curve shape for initial values of rowInc.
- ☐ However, the time taken for matrix multiplication roughly follows the order as:

$$\text{Bounded CAS} \geq \text{CAS} \geq \text{TAS} \geq \text{Atomic}$$

### Experiment 3:

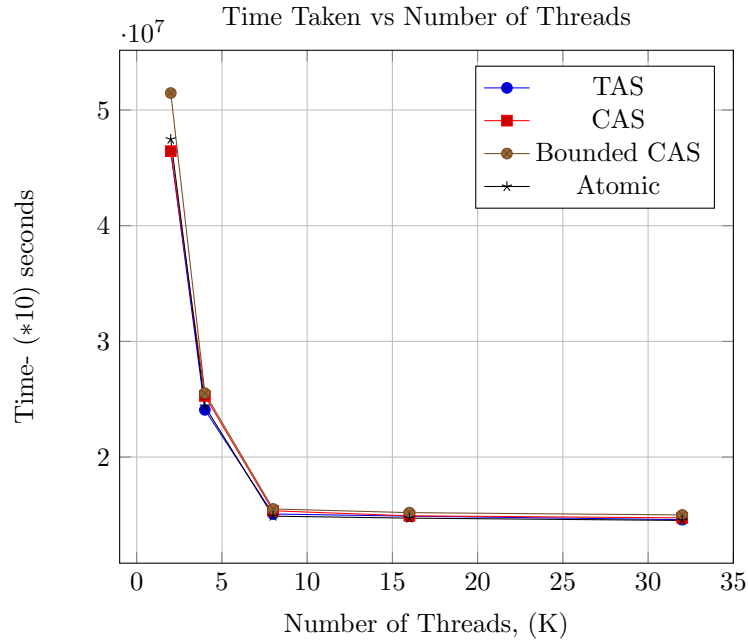
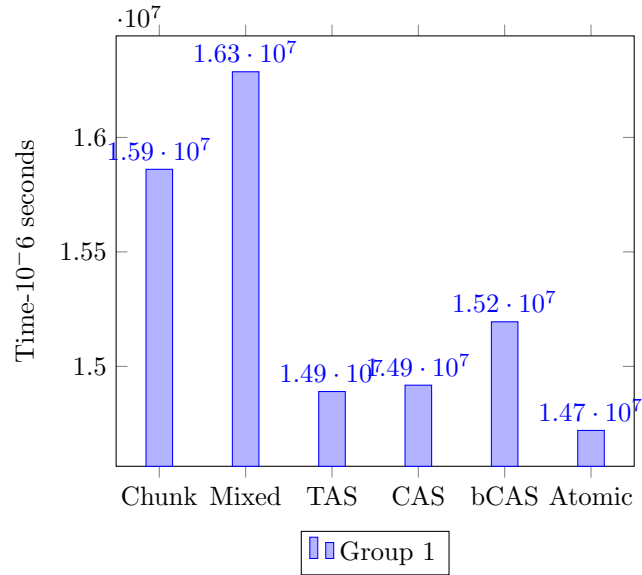


Figure 3: Time vs Number of Threads (K)

#### Observations:

- ☐ As the size of the matrix, N increases, the time taken to calculate the square of the matrix also increases.
- ☐ All the algorithms, i.e. TAS, CAS, Bounded CAS and Atomic methods takes almost the same time.
- ☐ Also the pattern in the curve as the N increases is almost same for all the curves.

## Experiment 4:



## Observations:

- Algorithms take time in the following order:

$$Mixed \geq Chunk \geq BoundedCAS \geq CAS \geq TAS \geq Atomic$$